Sparse Gaussian

January 3, 2017

1 mathematical notations

We first give brief description of mathematical notations will be used through out the project.

The original data set will be denoted as $\mathcal{D}$ which consists of N d-dimensional vectors $\boldsymbol{X} = \{\boldsymbol{x}^{(i)} = (x_1, \ldots, x_d) \,|\, i = 1, \ldots, N\}$. Let the new input data be $\boldsymbol{x}^* = (x_1^*, \ldots, x_d^*)$. The pseudo input data set is denoted as $\bar{\mathcal{D}}$ consists of $\bar{\boldsymbol{X}} = \{\bar{\boldsymbol{x}}^{(i)} = (x_1, \ldots, x_d) \,|\, i = 1, \ldots, M\}$. $\boldsymbol{X}$ is paired with target $\boldsymbol{Y} = (y^{(1)}, \ldots, y^{(N)})$, notice that $y^{(i)}$ are scalars. $\boldsymbol{x}^*$ is paired with new target y^*. The underlining latent function is denoted as $\boldsymbol{f}(\boldsymbol{x}) = \boldsymbol{y}$ and the pseudo one is $\bar{\boldsymbol{f}}$. A Gaussian distribution is denoted as $\mathcal{N}(\boldsymbol{f}|\boldsymbol{m}, \boldsymbol{V})$ with mean $\boldsymbol{m}$ and variance $\boldsymbol{V}$.

2 sparse Gaussian process

We first give a zero mean Gaussian prior over the underlining latent function: $p(\boldsymbol{f}|\boldsymbol{X}) = \mathcal{N}(\boldsymbol{f}|\boldsymbol{0}, \boldsymbol{K}_N)$ where $\boldsymbol{K}_N$ is our kernel matrix with elements given by, $[\boldsymbol{K}_N]_{ij} \equiv K_{\boldsymbol{x}^{(i)}\boldsymbol{x}^{(j)}} = K(\boldsymbol{x}^{(i)}, \boldsymbol{x}^{(j)})$: Notice that this is the case that we have same number of $\boldsymbol{x}^{(i)}, \boldsymbol{x}^{(j)}$. In case of different sizes, we use $\boldsymbol{K}_{NM}$, i.e. N rows for the first input matrix, M rows for the second input matrix.

$$K(\boldsymbol{x}^{(i)}, \boldsymbol{x}^{(j)}) = c \exp[-\frac{1}{2} \sum_{k=1}^{D} b_k (x_k^{(i)} - x_k^{(j)})^2], \quad \boldsymbol{\theta} \equiv \{c, \boldsymbol{b}\}, \tag{1}$$

where $\boldsymbol{\theta}$ is the hyperparameters. We provide noises to $\boldsymbol{f}$ such that $p(\boldsymbol{y}|\boldsymbol{f}) = \mathcal{N}(\boldsymbol{y}|\boldsymbol{f}, \sigma^2 \boldsymbol{I})$. By integrating out the latent function we have the marginal likelihood

$$p(\boldsymbol{y}|\boldsymbol{X}, \boldsymbol{\theta}) = \mathcal{N}(\boldsymbol{y}|\boldsymbol{0}, \boldsymbol{K}_N + \sigma^2 \boldsymbol{I}) \tag{2}$$

For prediction, the new input $\boldsymbol{x}^*$ conditioning on the observed data and hyperparameters. Let write the joint probability first

$$p(y^*, \boldsymbol{y}|\boldsymbol{x}^*, \mathcal{D}, \boldsymbol{\theta}) = \mathcal{N}\left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{pmatrix} K_{\boldsymbol{x}^*\boldsymbol{x}^*} + \sigma^2 & \boldsymbol{K}_{\boldsymbol{x}^*N} \\ \boldsymbol{K}_{\boldsymbol{x}^*N}^T & \boldsymbol{K}_N + \sigma^2 \boldsymbol{I} \end{pmatrix}\right), \tag{3}$$

where $\boldsymbol{K}_{\boldsymbol{x}^*N} = (K(\boldsymbol{x}^*, \boldsymbol{x}^{(1)}), \ldots, K(\boldsymbol{x}^*, \boldsymbol{x}^{(N)}))$, i.e. $[\boldsymbol{K}_{\boldsymbol{x}^*N}]_i = K(\boldsymbol{x}^*, \boldsymbol{x}^{(i)})$, and $K_{\boldsymbol{x}^*\boldsymbol{x}^*} = K(\boldsymbol{x}^*, \boldsymbol{x}^*)$. Now we can condition on $\boldsymbol{y}$ and get

$$\begin{aligned} &p(y^*|\boldsymbol{y}, \boldsymbol{x}^*, \mathcal{D}, \boldsymbol{\theta}) \\ &= \mathcal{N}(y^*|\boldsymbol{K}_{\boldsymbol{x}^*N}^T(\boldsymbol{K}_N + \sigma^2 \boldsymbol{I})^{-1}\boldsymbol{y}, K_{\boldsymbol{x}^*\boldsymbol{x}^*} + \sigma^2 - \boldsymbol{K}_{\boldsymbol{x}^*N}^T(\boldsymbol{K}_N + \sigma^2 \boldsymbol{I})^{-1}\boldsymbol{K}_{\boldsymbol{x}^*N}). \end{aligned} \tag{4}$$

For detailed proof, check Theorem 4.3.1 in Murphy's machine learning a probabilistic perspective.

Now we consider pseudo input $\bar{\boldsymbol{X}}$. Everything still holds except that there are no noises in it. The new input and target pair $(\boldsymbol{x}^*, y^*)$ is replaced by one of the actually data set and targets pairs $(\boldsymbol{x}^{(i)}, y_i)$. We therefore just use $\bar{\boldsymbol{f}}$ represents the pseudo outputs and $\bar{\boldsymbol{\theta}}$, and the single point likelihood is given by

$$p(y|\boldsymbol{x}, \bar{\boldsymbol{f}}, \bar{\boldsymbol{X}}) = \mathcal{N}(y|\boldsymbol{K}_{\boldsymbol{x}M}^T \boldsymbol{K}_M^{-1} \bar{\boldsymbol{f}}, K_{\boldsymbol{x}\boldsymbol{x}} + \sigma^2 - \boldsymbol{K}_{\boldsymbol{x}M}^T \boldsymbol{K}_M^{-1} \boldsymbol{K}_{\boldsymbol{x}M}), \tag{5}$$

where $\boldsymbol{K}_{\boldsymbol{x}M} = (K(\boldsymbol{x}, \bar{\boldsymbol{x}}^{(1)}), \ldots, K(\boldsymbol{x}, \bar{\boldsymbol{x}}^{(M)}))$, i.e. $[\boldsymbol{K}_{\boldsymbol{x}M}]_i = K(\boldsymbol{x}, \bar{\boldsymbol{x}}^{(i)})$. As the target data are i.i.d given the inputs, the complete data likelihood is given by

$$p(\boldsymbol{y}|\boldsymbol{X}, \bar{\boldsymbol{f}}, \bar{\boldsymbol{X}}) = \prod_{i=1}^{N} p(y_i|\boldsymbol{x}^{(i)}, \bar{\boldsymbol{f}}, \bar{\boldsymbol{X}}) = \mathcal{N}(\boldsymbol{y}|\boldsymbol{K}_{NM} \boldsymbol{K}_M^{-1} \bar{\boldsymbol{f}}, \boldsymbol{\Lambda} + \sigma^2 \boldsymbol{I}), \tag{6}$$

where $\boldsymbol{\Lambda} = \mathrm{diag}(\boldsymbol{\lambda}), \lambda_i = K_{\boldsymbol{x}^{(i)}\boldsymbol{x}^{(i)}} - \boldsymbol{K}_{\boldsymbol{x}^{(i)}M}^T \boldsymbol{K}_M^{-1} \boldsymbol{K}_{\boldsymbol{x}^{(i)}M}$, is a $N \times N$ diagonal matrix, and $[\boldsymbol{K}_{NM}]_{ij} = K(\boldsymbol{x}^{(i)}, \boldsymbol{x}^{(j)})$. Together with a Gaussian prior, $p(\bar{\boldsymbol{f}}|\bar{\boldsymbol{X}}) = \mathcal{N}(\bar{\boldsymbol{f}}|\boldsymbol{0}, \boldsymbol{K}_M)$, integrate over Eq.6 we have the SPGP marginal likelihood over pseudo inputs

$$\begin{aligned} p(\boldsymbol{y}|\boldsymbol{X}, \bar{\boldsymbol{X}}) &= \int p(\boldsymbol{y}|\boldsymbol{X}, \bar{\boldsymbol{f}}, \bar{\boldsymbol{X}}) p(\bar{\boldsymbol{f}}|\bar{\boldsymbol{X}}) \, \mathrm{d}\bar{\boldsymbol{f}} \\ &= \mathcal{N}(\boldsymbol{y}|\boldsymbol{0}, \boldsymbol{K}_{NM} \boldsymbol{K}_M^{-1} \boldsymbol{K}_{MN} + \boldsymbol{\Lambda} + \sigma^2 \boldsymbol{I}). \end{aligned} \tag{7}$$

Same as we have done from Eq.3 to Eq.4, we first write the joint probability of $y^*, \boldsymbol{y}$

$$\begin{aligned} &p(y^*, \boldsymbol{y}|\boldsymbol{x}^*, \boldsymbol{X}, \bar{\boldsymbol{X}}) \\ &= \mathcal{N}\left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{pmatrix} K_{\boldsymbol{x}^*\boldsymbol{x}^*} + \sigma^2 & \boldsymbol{K}_{\boldsymbol{x}^*M} \boldsymbol{K}_M^{-1} \boldsymbol{K}_{\boldsymbol{x}^*N} \\ (\boldsymbol{K}_{\boldsymbol{x}^*M} \boldsymbol{K}_M^{-1} \boldsymbol{K}_{\boldsymbol{x}^*N})^T & \boldsymbol{K}_{NM} \boldsymbol{K}_M^{-1} \boldsymbol{K}_{MN} + \boldsymbol{\Lambda} + \sigma^2 \boldsymbol{I} \end{pmatrix} \right), \end{aligned} \tag{8}$$

where $\boldsymbol{K}_{\boldsymbol{x}^*M} = (K(\boldsymbol{x}^*, \bar{\boldsymbol{x}}^{(1)}), \ldots, K(\boldsymbol{x}^*, \bar{\boldsymbol{x}}^{(M)}))$, i.e. $[\boldsymbol{K}_{\boldsymbol{x}^*M}]_i = K(\boldsymbol{x}^*, \bar{\boldsymbol{x}}^{(i)})$. From now on we let

$$\boldsymbol{Q}_{\boldsymbol{X},\boldsymbol{X}'} \equiv \boldsymbol{Q}(\boldsymbol{X}, \boldsymbol{X}') = \boldsymbol{K}_{\boldsymbol{X}M} \boldsymbol{K}_M^{-1} \boldsymbol{K}_{M\boldsymbol{X}'} \tag{9}$$

$$\boldsymbol{Q}_N = \boldsymbol{K}_{NM} \boldsymbol{K}_M^{-1} \boldsymbol{K}_{MN}, \tag{10}$$

Also, remember that here N and M represents input and pseudo input data set, matrices, as input matrices of $\boldsymbol{K}$, respectively. And after conditioning on $\boldsymbol{y}$, we have the SPGP predictive distribution

$$p(y^*|\boldsymbol{y}, \boldsymbol{x}^*, \boldsymbol{X}, \bar{\boldsymbol{X}}) = \mathcal{N}(\mu^*, {\sigma^*}^2) \tag{11}$$

$$\begin{aligned} \mu^* &= \boldsymbol{Q}_{\boldsymbol{x}^*N} (\boldsymbol{Q}_N + \boldsymbol{\Lambda} + \sigma^2 \boldsymbol{I})^{-1} \boldsymbol{y} \\ {\sigma^*}^2 &= K_{\boldsymbol{x}^*\boldsymbol{x}^*} - \boldsymbol{Q}_{\boldsymbol{x}^*N} (\boldsymbol{Q}_N + \boldsymbol{\Lambda} + \sigma^2 \boldsymbol{I})^{-1} \boldsymbol{Q}_{N\boldsymbol{x}^*} + \sigma^2 \end{aligned} \tag{12}$$